

Slide 1

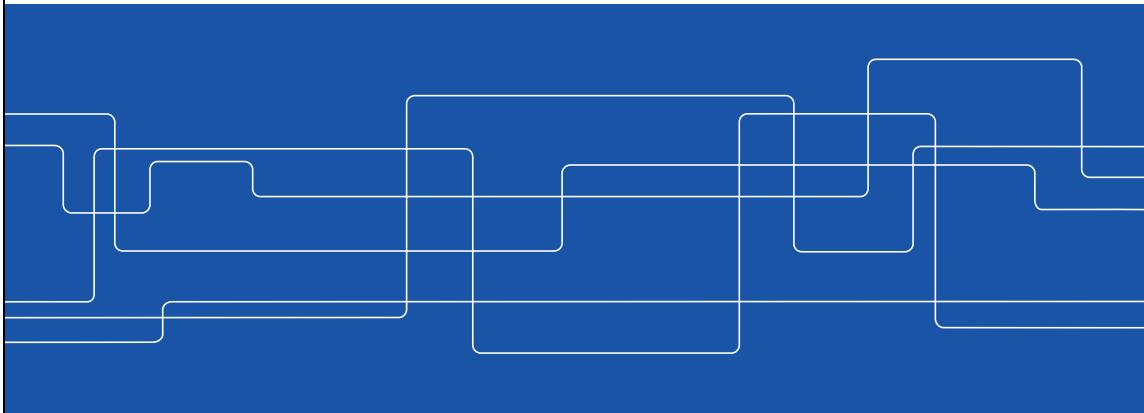


IK1552

Internetworking/Internetteknik

prof. Gerald Q. Maguire Jr. <http://people.kth.se/~maguire/>

School of Electrical Engineering and Computer Science (EECS), KTH Royal Institute of Technology
IK1552 Spring 2019, Period 4 2019.03.13 © 2019 G. Q. Maguire Jr. All rights reserved.



Slide 2



Module 4: UDP and friends

Lecture notes of G. Q. Maguire Jr.

For use in conjunction with James F. Kurose and Keith W. Ross, *Computer Networking: A Top-Down Approach*.

IK1552

SPRING 2019

SLIDE 2

Slide 3



Outline

- UDP
- Socket API
- BOOTP
- DHCP
- DNS, DDNS

Slide 4



Transport layer protocols

The transport layer is responsible for end-to-end delivery of entire messages

- uses Protocol and/or Port Number to demultiplex the incoming packet so that it can be delivered to a specific process
- Segmentation and Reassembly
 - Divides a message into transmittable segments and reassemble them at the receive
- Connection control - for connection-oriented transport protocols
- End-to-end **Flow** Control (in contrast to link level flow control)
- End-to-end **Error** Control (in contrast to link level error control)

Slide 5



Main Transport layer protocols

Three main transport layer protocols:

- User Datagram Protocol (UDP) <<< today's topic
Connectionless **unreliable** service
- Transmission Control Protocol (TCP)
Connection-oriented **reliable** stream service
- Stream Control Transmission Protocol (STCP)
a modern transmission protocol with many facilities
which the user can chose from

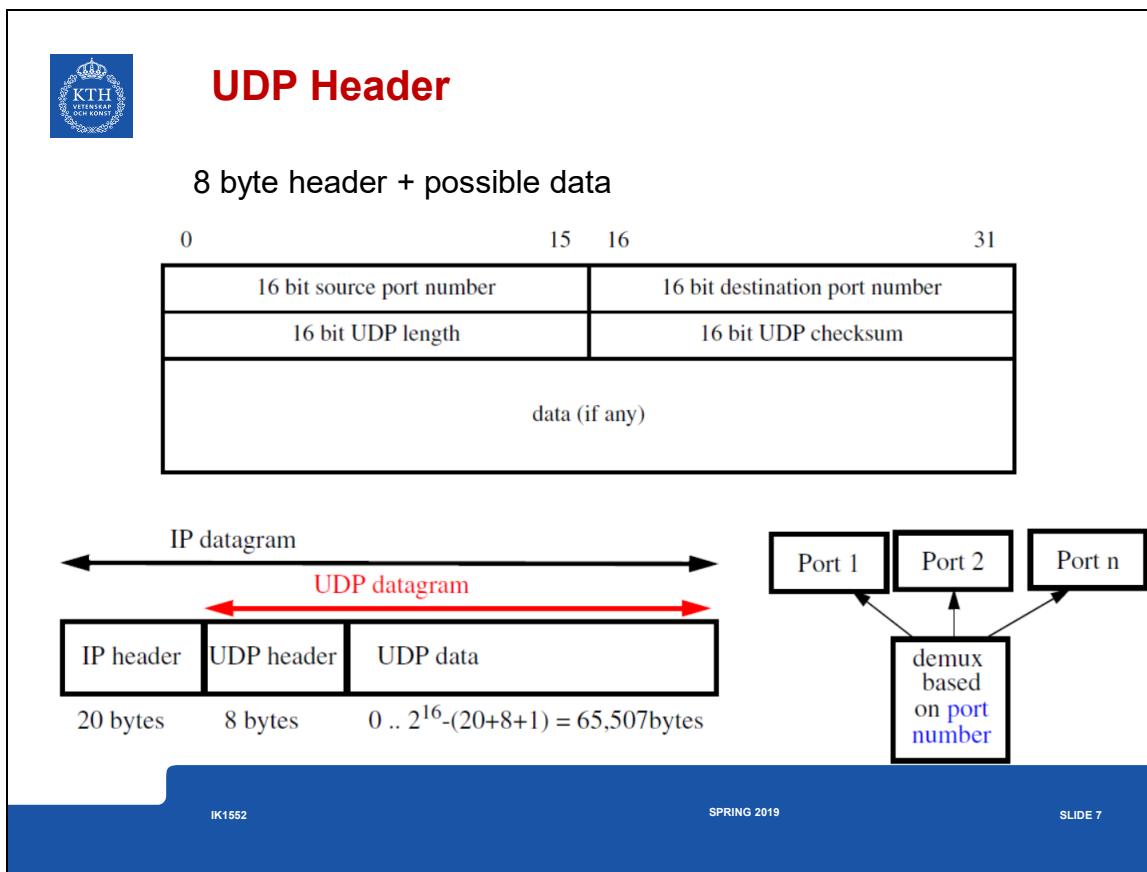
Slide 6



User Datagram Protocol (UDP)

- Datagram-oriented transport layer protocol
- Provides connectionless unreliable service
- No reliability guarantee
- Checksum covers both header and data, end-to-end, but optional
 - if you care about your data you should be doing end-to-end checksums or using an even stronger error detection (e.g., MD5).
- An UDP datagram is silently discarded if checksum is in error.
 - No error message is generated
- Lots of UDP traffic is only sent locally
 - thus the reliability is comparable to the error rate on the local links. (see Stevens, Vol. 1, figure 11.5, pg. 147 for comparison of Ethernet, IP, UDP, and TCP checksum errors)
- Each output operation results in **one** UDP datagram, which causes **one** IP datagram to be sent
- Applications which use UDP: DNS, TFTP, BOOTP, DHCP, SNMP, NFS, VoIP, etc.
 - An advantage of UDP is that it is a base to build your own protocols on.
 - Especially if you do not need reliability and in order delivery of lots of data.

Slide 7



Slide 8

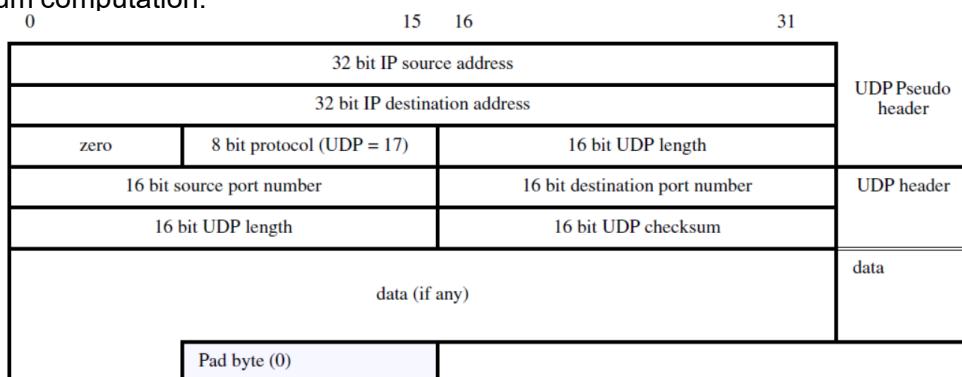


UDP Checksum and Pseudo-Header

UDP checksum covers more info than is present in the UDP datagram alone:
pseudo-header and **pad byte (0)** {to even number of 16 bit words}.

Propose: to verify the UDP datagram reached its correct destination: **right port number at the right IP address**.

Pseudo-header and pad byte are **not transmitted** with the UDP datagram, only used for checksum computation.



Slide 9



Reserved and Available UDP Port Numbers

	keyword	UNIX keyword	Description
0	reserved		
7	ECHO	echo	Echo
9	DISCARD	discard	Discard == sink null
11	USERS	systat	Active users
13	DAYTIME	daytime	Daytime
15	-	netstat	Network status program
17	QUOTE	qtd	Quote of the day
19	CHARGEN	chargen	Character generator
37	TIME	time	Time server
39	RLP	rip	Resource Location Protocol
42	NAMESERVER	name	Host Name Server
43	NICNAME	whois	Who is
53	DOMAIN	domain	Domain Name Server
67	BOOTPS	bootps	Bootstrap Protocol Server
68	BOOTPC	bootpc	Bootstrap Protocol Client
69	TFTP	tftp	Trivial File Transfer Protocol
88	KERBEROS	kerberos5	Kerberos v5 kdc
111	SUNRPC	sunrpc	SUN Remote Procedure Call (portmap)
123	NTP	ntp	Network Time Protocol
137	netbios_ns	netbios_ns	NetBIOS name service
138	netbios_dgm	netbios_dgm	NetBIOS Datagram Service
139	netbios_ssn	netbios_ssn	NetBIOS Session Service
161	snmp	Simple Network Management Protocol Agent	
162	snmp-trap	Simple Network Management Protocol Traps	
512	biff	mail notification	
513	who	remote who and uptime	
514	syslog	remote system logging	
517	talk	conversation	
520	route	routing information protocol	
525	timed	remote clock synchronization	
533	netwall	netwall	Emergency broadcasting
750	kerberos	kerberos	Kerberos (server)
6000 + display number		X11 server	
7000	X11 font server		

Slide 10



Port numbers in three groups

Range	Purpose
0 .. 1023	System (Well-known) Ports [†]
1024 .. 49151	User (Registered) Ports
49152 .. 65535	Dynamic and/or Private Ports

[†]Roughly 300 well known port numbers remain unassigned and 38 reserved

Roughly 26k registered port numbers remain unassigned and 9 reserved

'For the purpose of providing services to **unknown** callers, a service contact port is defined. This list specifies the port used by the server process as its contact port. The contact port is sometimes called the "well-known port".'

<http://www.iana.org/assignments/port-numbers>

Linux chooses the local port to use for TCP and UDP traffic from this range:

```
$ cat /proc/sys/net/ipv4/ip_local_port_range  
32768 61000
```

Slide 11

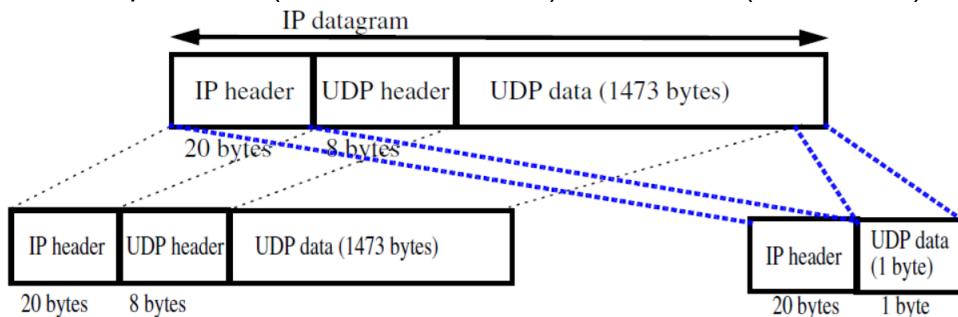


MTU and Datagram Fragmentation

If datagram size > MTU, perform fragmentation.

- At sending host or at intermediate router (IPv4).
- Reassembled only at final destination.

Example: 1501 (20 + 8 + 1473 data) on Ethernet (MTU=1500):



Note there is no UDP header in the second fragment.

Therefore, a frequent operation is to compute the path MTU before sending anything else. (see RFC 1191 for the table of common MTUs)

J. C. Mogul and S. E. Deering, 'Path MTU discovery', *Internet Request for Comments*, vol. RFC 1191 (Draft Standard), Nov. 1990 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1191.txt>

J. McCann, S. Deering, and J. Mogul, 'Path MTU Discovery for IP version 6', *Internet Request for Comments*, vol. RFC 1981 (Draft Standard), Aug. 1996 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1981.txt>

B. Black and K. Kompella, 'Maximum Transmission Unit Signalling Extensions for the Label Distribution Protocol', *Internet Request for Comments*, vol. RFC 3988 (Experimental), Jan. 2005 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc3988.txt>

M. Mathis and J. Heffner, 'Packetization Layer Path MTU Discovery', *Internet Request for Comments*, vol. RFC 4821 (Proposed Standard), Mar. 2007 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc4821.txt>

Slide 12



Fragmentation Required

If (datagram size > MTU) and DF (Don't Fragment) in IP header is on, then the router sends ICMP Unreachable Error.
⇒ this can be used to find Path MTU.

Slide 13



Interaction between UDP and ARP

With ARP cache empty, send a UDP datagram with 8192 bytes onto an Ethernet

- 8192 bytes > Ethernet MTU, therefore 6 fragments are created by IP
- if ARP cache is empty, first fragment causes ARP request to be sent

This leads to two timing questions:

1. Are the remaining fragments sent before the ARP reply is received?
2. What does ARP do with multiple packets to the same destination while waiting for a reply?

Example under BSD

```
Bsd% arp -a ARP cache is empty
Bsd% sock -u -i -n1 -w8192 svr4 discard
 1  0.0      arp who-has svr4 tell bsdi
 2  0.001234  (0.0012)  arp who-has svr4 tell bsdi
 3  0.001941  (0.0007)  arp who-has svr4 tell bsdi
 4  0.002775  (0.0008)  arp who-has svr4 tell bsdi
 5  0.003495  (0.0007)  arp who-has svr4 tell bsdi
 6  0.004319  (0.0008)  arp who-has svr4 tell bsdi
 7  0.008772  (0.0045)  arp reply svr4 is-at 0:0:c0:c2:9b:26
 8  0.009911  (0.0011)  arp reply svr4 is-at 0:0:c0:c2:9b:26
 9  0.011127  (0.0012)  bsdi > svr4: (frag 10863:800@7400)
10  0.01255   (0.0001)  arp reply svr4 is-at 0:0:c0:c2:9b:26
11  0.012562  (0.0013)  arp reply svr4 is-at 0:0:c0:c2:9b:26
12  0.013458  (0.0009)  arp reply svr4 is-at 0:0:c0:c2:9b:26
13  0.014526  (0.0011)  arp reply svr4 is-at 0:0:c0:c2:9b:26
14  0.015583  (0.0011)  arp reply svr4 is-at 0:0:c0:c2:9b:26
```

Slide 14



On a BSDI system

- each of the additional (5) fragments caused an ARP request to be generated
this violates the Host Requirements RFC - which tries to prevent **ARP flooding** by limiting the maximum rate to 1 per second
- when the ARP reply is received the last fragment is sent
Host Requirements RFC says that ARP should save at least one packet and this should be the latest packet
- unexplained anomaly: the System Vr4 system sent 7 ARP replies back!
- no ICMP “time exceeded during reassembly” message is sent
 - BSD derived systems - never generate this error! It does set the timer internally and discard the fragments, but never sends an ICMP error.
 - fragment 0 (which contains the UDP header) was not received - so there is no way to know which process sent the fragment; thus unless fragment 0 is received - you are not required to send an ICMP “time exceeded during reassembly” error.

Not just a fluke (i.e., a rare event)

- The same error occurs even if you don't have fragmentation - simply sending multiple UDP datagrams rapidly when there is no ARP entry is sufficient!
- NFS sends UDP datagrams whose length just exceeds 8192 bytes
- NFS will timeout and resend
- however, there will always be this behavior - if the ARP cache has no entry for this destination!

Slide 15



Still a problem?

UDP with 8192 payload to echo port as seen on SuSE 9.2 linux 2.6.8-24:

No.	Time	Source	Destination	Protocol	Info
37	3.020002	172.16.33.16	Broadcast	ARP	Who has 172.16.33.5? Tell 172.16.33.16
38	3.021385	172.16.33.5	172.16.33.16	ARP	172.16.33.5 is at 00:40:8c:24:37:f4
39	3.021422	172.16.33.16	172.16.33.5	IP	Fragmented IP protocol (proto=UDP 0x11, off=4440)
40	3.021452	172.16.33.16	172.16.33.5	IP	Fragmented IP protocol (proto=UDP 0x11, off=5920)
41	3.021480	172.16.33.16	172.16.33.5	IP	Fragmented IP protocol (proto=UDP 0x11, off=7400)

3.021385-3.020002=.001383 sec. \Rightarrow 1.383ms for the ARP reply

All but the last 3 fragments are dropped! Including the initial echo request packet -- so in the fragments that do arrive you don't know who they are for -- because the first fragment was lost!

Slide 16



With an even larger UDP packet

I removed the arp cache entry with: /sbin/arp -i eth1 -d 172.16.33.5

When sending 65500 bytes of UDP payload -- it loses many packets (in fact all but the last 3 fragments)!!!

No. Time Source Destination Protocol Info

```
36 4.342158 172.16.33.16 Broadcast ARP Who has 172.16.33.5? Tell 172.16.33.16
37 4.342875 172.16.33.5 172.16.33.16 ARP 172.16.33.5 is at 00:40:8c:24:37:f4
38 4.342906 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=62160)
39 4.342932 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=63640)
40 4.342986 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=65120)
```

With the entry in the ARP cache get:

No. Time Source Destination Protocol Info

```
35 5.118063 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=1480)
36 5.118095 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=2960)
37 5.118115 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=4440)
38 5.118214 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=5920)
39 5.118328 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=7400)
...
75 5.122787 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=60680)
76 5.122877 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=62160)
77 5.122999 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=63640)
78 5.123122 172.16.33.16 172.16.33.5 IP Fragmented IP protocol (proto=UDP 0x11, off=65120)
```

The initial UDP Echo request is still lost! The key parameter is

/proc/sys/net/ipv4/neigh/ethX/unres_qlen where X is the interface (i.e., eth0, eth1, ...)

-- the default value is 3 (In 2016, in OpenSuse 13.2 the value is 31)

Slide 17



Maximum UDP Datagram size

theoretical limit: 65,535 bytes - due to (IP's) 16-bit total length field

with 20 bytes of IP header + 8 bytes of UDP header \Rightarrow 65,507 bytes of user data

two limits:

- sockets API limits size of send and receive buffer; generally 8 kbytes, but you can call a routine to change this (modern linux kernels auto tune this buffer size)
- TCP/IP implementation - Stevens found various limits to the sizes - even with loopback interface (see Stevens, Vol. 1, pg. 159)

Hosts are required to handle at least 576 byte IP datagrams, thus lots of protocols limit themselves to 512 bytes or less of data:

DNS, TFTP, BOOTP, and SNMP

Slide 18



Datagram truncation

What if the application is not prepared to read the datagram of the size sent?

Implementation dependent:

- traditional Berkeley: silently truncate
- 4.3BSD and Reno: **can** notify the application that the data was truncated
- SVR4: excess data returned in subsequent reads - application is not told that this all comes from one datagram
- TLI: sets a flag that more data is available, subsequent reads return the rest of the datagram

Slide 19



Socket API

```
int socket(int domain, int type, int protocol);
    creates an endpoint description that you can use to send and receive network traffic
int bind(int sockfd, struct sockaddr *my_addr, socklen_t addrlen);
    binds a socket to the local address and port: my_address
int connect(int sockfd, const struct sockaddr *serv_addr, socklen_t addrlen);
    connect the local socket to a remote socket
int listen(int s, int backlog);
    indicates socket is willing to accept connections & limits the queue of pending connections
int accept(int s, struct sockaddr *addr, socklen_t *addrlen);
    accepted the first connection request from the queue and connects it to the socket
connection oriented sending and receiving:
    ssize_t send(int s, const void *buf, size_t len, int flags);
    ssize_t recv(int s, void *buf, size_t len, int flags);
datagram sending and receiving:
    ssize_t sendto(int s, const void *buf, size_t len, int flags, const struct sockaddr *to, socklen_t tolen);
    ssize_t recvfrom(int s, void *buf, size_t len, int flags, struct sockaddr *from, socklen_t *fromlen);
int close(int fd) and int shutdown(int s, int how) - end it all!
```

Slide 20



Windows Socket API

Have to start up the Windows TCP/IP stack:

```
int WSASStartup( _In_ WORD wVersionRequested,  
                  _Out_ LPWSADATA lpWSAData);
```

To cleanly shutdown the stack:

```
WSACleanup();
```

There are additional WSA functions, for example for asynchronous I/O

Slide 21



Learning about Socket programming

For examples of using the socket API and networking programming see :

- Richard Stevens, *UNIX Network Programming, Volume 1, Second Edition: Networking APIs: Sockets and XTI*, Prentice Hall, 1998, ISBN 0-13-490012-X
<http://www.kohala.com/start/unpv12e.html>
- Brian "Beej" Hall, "Beej's Guide to Network Programming: Using Internet Sockets", Jorgensen Publishing, October 21, 2011, 142 pages, ASIN: B002AD9SNK <http://beej.us/guide/bgnet/>

Two addition socket functions for controlling various properties of sockets are:

- int `getsockopt`(int s, int level, int optname, void *optval, socklen_t *optlen);
- int `setsockopt`(int s, int level, int optname, const void *optval, socklen_t optlen);

For example, the options SO_SNDBUF and SO_RCVBUF - control the size of the sending buffer and the receiver buffer.

Brian 'Beej' Hall, *Beej's Guide to Network Programming: Using Internet Sockets*. Jorgensen Publishing, 2011, ISBN: ASIN: B002AD9SNK [Online]. Available: <http://beej.us/guide/bgnet/>

Slide 22



Simple UDP client

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

#define bigBufferSize 8192
#define destination_host "127.0.0.1"
#define Destination_Port 52000

main(argc, argv)
int argc;
char **argv;
{ int client_socket_fd; /* Socket to client, server */
  struct sockaddr_in server_addr; /* server's address */
  char bigBuffer[bigBufferSize];
  int sendto_flags=0;

  /* create a UDP socket */
  if ((client_socket_fd = socket(AF_INET, SOCK_DGRAM,
                                 IPPROTO_UDP)) == -1)
    { perror("Unable to open socket"); exit(1); }

  /* initialize the server address structure */
  memset((char*)&server_addr, 0, sizeof(server_addr));
  server_addr.sin_family=AF_INET;
  server_addr.sin_port=htons(Destination_Port);
  if (inet_aton(destination_host, (struct sockaddr*)&server_addr.sin_addr) == 0)
    { fprintf(stderr, "could not get an address for: %s", destination_host); exit(1); }

  sprintf(bigBuffer, "This is a simple test string to be sent to the other party\n");

  if ((sendto(client_socket_fd,
              bigBuffer, strlen(bigBuffer),
              sendto_flags, (struct sockaddr*)&server_addr,
              sizeof(server_addr))) == -1)
    { perror("Unable to send to socket"); close(client_socket_fd); exit(1); }

  close(client_socket_fd); /* close the socket */
  exit(0);
}
```

IK1552

SPRING 2019

SLIDE 22

Slide 23



UDP server design

Stevens, Vol, 1, pp. 162-167 discusses how to program a UDP server

You can often determine what IP address the request was sent to (i.e., the destination address):

for example: thus ignoring datagrams sent to a broadcast address

You can limit a server to a given incoming IP address:

thus limiting requests to a given interface

You can limit a server to a given foreign IP address and port:

only accepting requests from a given foreign IP address and port #

Multiple recipients per port (for implementations with multicasting support)

setting SO_REUSEADDR socket option ⇒ each process gets a copy of the incoming datagram

Note: limited size input queue to each UDP port, can result in silent discards without an ICMP message being sent back (since **OS discarded**, not the network!)

Slide 24



UDP listener example

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <arpa/inet.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

#define bigBufferSize 8192
#define my_port      52000

main(argc, argv)
int argc;
char *argv;
{
    int client_socket_fd; /* Socket to client, server */
    struct sockaddr_in client_addr; /* client's address */
    struct sockaddr_in other_addr; /* other party's address */
    socklen_t other_addr_len;
    char source_host_address_as_string[INET_ADDRSTRLEN+2];

    char bigBuffer[bigBufferSize];
    int sendto_flags=0;

    ssize_t nbytes;
    other_addr_len = sizeof(other_addr);

    /* create a UDP socket */
    if ((client_socket_fd = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP)) == -1) {
        perror("Unable to open socket");
        exit(1);
    }

    /* initialize the server address structure */
    memset((char*)&client_addr, 0, sizeof(client_addr));
    client_addr.sin_family=AF_INET;
    client_addr.sin_port=htons(my_port);
    client_addr.sin_addr.s_addr = htonl(INADDR_ANY);

    if (bind(client_socket_fd,(struct sockaddr*)&client_addr,sizeof(client_addr))== -1)
        { close(client_socket_fd); exit(1); }

    if ((nbytes = recvfrom(client_socket_fd,bigBuffer,bigBufferSize, sendto_flags,
                          (struct sockaddr *)&other_addr, &other_addr_len)) == -1)
        { perror("Unable to send to socket"); close(client_socket_fd); exit(1); }

    /* two different ways of printing the IP address that send the UDP datagram */
    if (inet_ntop(AF_INET,&(other_addr.sin_addr),source_host_address_as_string,
                  INET_ADDRSTRLEN))
    {
        printf("Received packet from %s:%d\n",
               source_host_address_as_string, ntohs(other_addr.sin_port));
    } else {
        printf("Received packet from %s:%d\n",
               inet_ntoa(other_addr.sin_addr), ntohs(other_addr.sin_port));
    }

    printf("Data (%d bytes): %s\nString length=%d\n", (int)nbytes, bigBuffer,
           strlen(bigBuffer));

    close(client_socket_fd); /* close the socket */
    exit(0);
}
```

IK1552

SPRING 2019

SLIDE 24

Slide 25



Building a UDP packet from scratch

```
/* simple example of building a UDP packet from scratch, based on the program: PingPong - 970621 by Willy TARREAU
<tarreau@aemliaif.ibp.fr>
http://www.insecure.org/sploits/inetd.internal_udp_ports.DOS.attack.html
As this program uses RAW sockets, you must be root to run it */
#include <netinet/in.h>
#include <netinet/ip.h>
#include <netinet/udp.h>
#include <netdb.h>
#include <stdlib.h>
#include <string.h>
#include <stdio.h>
#include <ctype.h>
#include <errno.h>
struct sockaddr_in addrfrom;
struct sockaddr_in addto;
int s;
u_char outpack[65536];
struct iphdr *ip;
struct udphdr *udp;
main(int argc, char **argv) {
    struct sockaddr_in *from;
    struct sockaddr_in *to;
    struct protoent *proto;
    int i;
    char *src,*dest;
    int srcp, destp;
    int packetsize,datasize;
    if (argc!=5) {fprintf(stderr,"Usage: %s src_addr src_port dst_addr dst_port\n", argv[0]);
        fprintf(stderr,"src_addr and dst_addr must be given as IP addresses (xxx.xxx.xxx.xxx)\n");
        exit(2);}
    src=argv[1]; srcp=atoi(argv[2]); dest=argv[3]; destp=atoi(argv[4]);
```

IK1552

SPRING 2019

SLIDE 25

Slide 26



```
if (!(proto = getprotobynumber("raw"))) {perror("getprotobynumber(raw)");exit(2);}

if ((s = socket(AF_INET, SOCK_RAW, proto->p_proto)) < 0) {perror("socket");exit(2);}
memset(&addrfrom, 0, sizeof(struct sockaddr));
from = (struct sockaddr_in *)&addrfrom;
from->sin_family = AF_INET;
from->sin_port=htons(srccp);
if (!inet_aton(src, &from->sin_addr)) {fprintf(stderr,"Incorrect address for 'from': %s\n",src);exit(2);}

memset(&addrto, 0, sizeof(struct sockaddr));
to = (struct sockaddr_in *)&addrto;
to->sin_family = AF_INET;
to->sin_port=htons(destp);
f (inet_aton(dest, &to->sin_addr)) {fprintf(stderr,"Incorrect address for 'to': %s\n",dest);exit(2);}

packetsize=0;                                /* build a UDP packet from scratch */
ip=(struct iphdr *)outpack;
ip->version=4;                                /* IPv4 */
ip->ihl=5;                                    /* IP header length: 5 words */
ip->tos=0;                                     /* no special type of service */
ip->id=0;                                      /* no ID */
ip->frag_off=0;                               /* not a fragment - so there is no offset */
ip->ttl=64;                                    /* TTL = 64 */
If (!(proto = getprotobynumber("udp"))) {perror("getprotobynumber(udp)"); exit(2);}
ip->protocol=proto->p_proto;
ip->check=0;                                    /* null checksum, will be automatically computed by the kernel */
ip->saddr=from->sin_addr.s_addr;
ip->daddr=to->sin_addr.s_addr;
/* end of ip header */
```

Slide 27



```
packetsize+=ip->ihl<<2;
                                         /* udp header */
udp=(struct udphdr *)((int)outpack + (int)(ip->ihl<<2));
udp->source=htons(srccp);
udp->dest=htons(destp);
udp->check=0;                                /* ignore UDP checksum */
packetsize+=sizeof(struct udphdr);           /* end of udp header */
                                         /* add data to UDP payload if you want: */

for (datasize=0;datasize<8;datasize++) {
    outpack[packetsize+datasize]='A'+datasize;
}
packetsize+=datasize;
udp->len=htons(sizeof(struct udphdr)+datasize);
ip->tot_len=htons(packetsize);

if (sendto(s, (char *)outpack, packetsize, 0, &addrto, sizeof(struct sockaddr))==-1)
    {perror("sendto"); exit(2);}
printf("packet sent !\n");
close(s);
exit(0);}
```

Slide 28

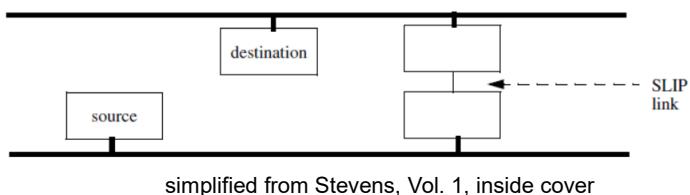


ICMP Source Quench Error

Since UDP has no flow control, a node could receive datagrams faster than it can process them. In this situation the host **may** send an ICMP source quench.

Note: “**may** be generated” - it is **not** required to generate this error

Stevens (Vol. 1, pp. 160-161) gives the example of sending 100 1024-byte datagrams from a machine on an Ethernet via a router and SLIP line to another machine:



- SLIP link is ~1000 times slower than the ethernet
- 26 datagrams are transmitted, then a source quench is sent for each successive datagram

Slide 29



ICMP Source Quench Error - continued

- the router gets all 100 packets, before the first has been sent across the link!
the Router Requirements RFC - says that routers should not generate source quench errors, since it just consumes network bandwidth and it is an ineffective and unfair fix for congestion
- In any case, the sending program never responded to the source quench errors!
 - BSD implementations ignore received source quenches if the protocol is UDP
 - the program finished before the source quench was received!

Thus if you want reliability you have to build it in and do end-to-end flow control, error checking, and use (and thus wait for) acknowledgements.

Slide 30



No error control

Since UDP has no error control, the sender has to take responsibility for sending the datagram again if this datagram must be delivered.

But how does the send know if a datagram was successfully delivered?

- Unless the receiver sends a reply (or does some action due to receiving a given datagram) the sender will **not** know!
[**loss**]
- Note that without some additional mechanism - the sender doesn't know if the datagram was delivered multiple times!
[**duplicates**]

If you want reliability you have to build your own protocol on top of UDP to achieve it. This includes deciding on your own retransmission scheme, timeouts, etc.

Slide 31



BOOTP: Bootstrap Protocol (RFC 951)

Although you can figure out who you are, i.e., your IP address, via RARP - many machines want more information.

BOOTP requests and answer are sent via UDP (port 67 server; port 68 client)

- so it is easy to make a user space server
- the client (who wants the answer) need not have a full TCP/IP, it can simply send what **looks** like a UDP datagram with a BOOTP request (see Stevens, Vol. 1, figure 16.2, pg. 2).

Opcode (1=request, 2=reply)	hardware type(1=ethernet)	hardware address length (6 for ethernet)	hop count		
transaction ID					
number of seconds		unused			
client IP address					
your IP address					
server IP address					
gateway IP address					
client hardware address (16 bytes of space)					
server hostname (64 bytes)					
Boot file name (128 bytes)					
Vendor specific information (64 bytes)					

W. J. Croft and J. Gilmore, ‘Bootstrap Protocol’, *Internet Request for Comments*, vol. RFC 951 (Draft Standard), Sep. 1985 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc951.txt>

Slide 32



BOOTP (continued)

When a request is sent as an IP datagram:

- if client does not know its IP address it uses 0.0.0.0
- if it does not know the server's address it uses 255.255.255.255
- if the client does not get a reply, it tries again in **about 2 sec.**

Slide 33



Vendor specific information (RFC 1497 and RFC1533)

if this area is used the first 4 bytes are: IP address 99.130.83.99 this is called the “**magic cookie**”, Then the rest of the area is a list of items, possibly including:

- Pad (tag=0);
- Subnet mask (tag=1);
- Time offset (tag=2);
- List of IP addresses of Gateways (tag=3);
- Time server’s IP address (tag=4);
- Name Server (tag=5);
- Domain Name Server (tag=6);
- LOG server (tag=7); ...
- LPR server (tag=9); ...
- this Host’s name (tag=12);
- Boot file size (tag=13); ...
- Domain name (tag=15); ...
- End (tag=255)

IK1552

SPRING 2019

SLIDE 33

W. J. Croft and J. Gilmore, ‘Bootstrap Protocol’, *Internet Request for Comments*, vol. RFC 951 (Draft Standard), Sep. 1985 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc951.txt>

P. A. Prindeville, ‘BOOTP vendor information extensions’, *Internet Request for Comments*, vol. RFC 1048, Feb. 1988 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1048.txt>

J. K. Reynolds, ‘BOOTP vendor information extensions’, *Internet Request for Comments*, vol. RFC 1084, Dec. 1988 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1084.txt>

J. Reynolds, ‘BOOTP Vendor Information Extensions’, *Internet Request for Comments*, vol. RFC 1395 (Draft Standard), Jan. 1993 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1395.txt>

J. Reynolds, ‘BOOTP Vendor Information Extensions’, *Internet Request for Comments*, vol. RFC 1497 (Draft Standard), Aug. 1993 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1497.txt>

S. Alexander and R. Droms, ‘DHCP Options and BOOTP Vendor Extensions’, *Internet Request for Comments*, vol. RFC 1533 (Proposed Standard), Oct. 1993 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1533.txt>

Slide 34



DHCP: Dynamic Host Configuration Protocol (RFC 1531)

Extends the Vendor specific options area by 312 bytes.

This protocol is designed to make it easier to allocate (and reallocate) addresses for clients. DHCP defines:

- **Requested IP Address** - used in client request (DHCPDISCOVER) to request that a particular IP address (tag=50)
- **IP Address Lease Time** - used in a client request (DHCPDISCOVER or DHCPREQUEST) to request a lease time for the IP address. In a server reply (DHCPOFFER), specific lease time offered. (tag=51)
- **Option Overload** - used to indicate that the DHCP “sname” or “file” fields are being overloaded by using them to carry DHCP options. A DHCP server inserts this option if the returned parameters will exceed the usual space allotted for options, i.e., it uses the sname and file fields for another purpose! (tag=52)

R. Droms, ‘Dynamic Host Configuration Protocol’, *Internet Request for Comments*, vol. RFC 1531 (Proposed Standard), Oct. 1993 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1531.txt>



DHCP (continued)

- **DHCP Message Type** - the type of the DHCP message (tag=53)
- **Server Identifier** - used in DHCPOFFER and DHCPREQUEST (optionally in DHCPACK and DHCPNAK) messages. Servers include this in the DHCPOFFER to allow the client to distinguish **between** lease offers. DHCP clients indicate which of several lease offers is being accepted by including this in a DHCPREQUEST message. (tag=54)
- **Parameter Request List** - used by a DHCP client to request values for specified configuration parameters. The client **may** list options in order of preference. The DHCP server **must** try to insert the requested options in the order requested by the client. (tag=55)

Slide 36



DHCP (continued)

- **Message** - used by a server to provide an error message to client in a DHCPNAK message in the event of a failure. A client may use this in a DHCPDECLINE message to indicate the reason **why** the client declined the offered parameters.(tag=56)
- **Maximum DHCP Message Size** - specifies the maximum length DHCP message that it is willing to accept. A client may use the maximum DHCP message size option in DHCPDISCOVER or DHCPREQUEST messages, but should not use the option in DHCPDECLINE messages. (tag=57)
- **Renewal (T1) Time Value** - specifies the time interval from address assignment until the client transitions to the RENEWING state. (tag=58)
- **Rebinding (T2) Time Value** - specifies the time interval from address assignment until the client transitions to the REBINDING state.(tag=59)

Slide 37



DHCP (continued)

- **Class-identifier** - used by DHCP clients to optionally identify the type and configuration of a DHCP client. Vendors and sites may choose to define specific class identifiers to convey particular configuration or other identification information about a client. Servers not equipped to interpret the class-specific information sent by a client **must** ignore it (although it may be reported). (tag=60)
- **Client-identifier** - used by DHCP clients to specify their unique identifier. DHCP servers use this value to index their database of address bindings. This value is expected to be unique for all clients in an administrative domain. (tag=61)

Slide 38



DHCP's importance

- allows reuse of address, which avoids having to tie up addresses for systems which are not currently connected to the Internet
- avoids user configuration of IP address (avoids mistakes and effort)
- allows recycling of an IP address when devices are scrapped
- ...

How big a problem is manual configuration?

A large site with > 70k IP addressable devices; or consider what happens if each of > 1.4 M Walmart employees has an IP device

Product	Vendor	URL
Network Registrar	Cisco	http://www.cisco.com
NetID	Nortel Networks	http://www.nortelnetworks.com
Meta IP	CheckPoint	http://www.checkpoint.com
VitalQIP	Alcatel-Lucent	http://enterprise.alcatel-lucent.com

Slide 39



DHCP performance problems

Most implementations of DHCP do a duplicate address detection (DAAD) test **after** they have picked an address to assign.

An alternative approach to speed up the DHCP process does the duplicate address detection process in the background (in advance) so that you will have a set of recently tested addresses to hand out:

Jon-Olov Vatn and Gerald Q. Maguire Jr., "The effect of using co-located care-of addresses on macro handover latency", Fourteenth Nordic Tele-traffic Seminar (NTS 14), August 18 - 20, 1998, Lyngby, Denmark.

<http://urn.kb.se/resolve?urn=urn:nbn:se:kth:diva-50201>

The result is that a DHCP request can be answered in less than 100ms.

Slide 40



Example of dhcpcd.conf

```
### Managed by Linuxconf, you may edit by hand.  
### Comments may not be fully preserved by linuxconf.  
server-identifier dhcptest1;  
default-lease-time 1000;  
max-lease-time 2000;  
option domain-name "3ctechnologies.se";  
option domain-name-servers 130.237.12.2;  
option host-name "s1.3ctechnologies.se";  
option routers 130.237.12.2;  
option subnet-mask 255.255.255.0;  
subnet 130.237.12.0 netmask 255.255.255.0 {  
    range 130.237.12.3 130.237.12.200;  
    default-lease-time 1000;  
    max-lease-time 2000;  
}  
subnet 130.237.11.0 netmask 255.255.255.0 {  
    range 130.237.11.3 130.237.11.254;  
    default-lease-time 1000;  
    max-lease-time 2000;  
}
```

Slide 41



DHCP and DNS

Interaction between DHCP and DNS is needed – see RFC 4702

For example: once a host is assigned an IP address the DNS should be updated dynamically:

- If the host has not got a name: it should assign a name along the IP address assignment (no DNS update is needed).
- If the host has already a name: the DNS should be dynamically updated once the host has gotten a new IP address from DHCP.

[RFC 4702] M. Stapp, B. Volz, and Y. Rekhter, ‘The Dynamic Host Configuration Protocol (DHCP) Client Fully Qualified Domain Name (FQDN) Option’, *Internet Request for Comments*, vol. RFC 4702 (Proposed Standard), Oct. 2006 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc4702.txt>



Trivial File Transfer Protocol (TFTP)

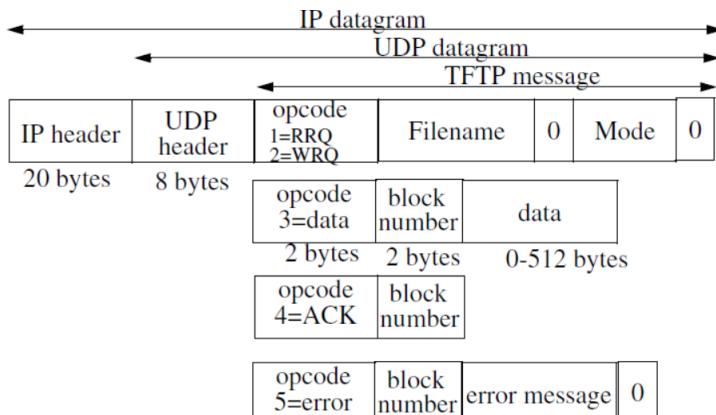
TFTP uses UDP (unlike FTP which uses TCP)

- simple and small
- requires only UDP, IP, and a device driver - easily fits in ROM
- a stop-and-wait protocol
- lost packets detected by timeout and retransmission
- Two operations:
 - Read Request (RRQ)
 - Write Request (WRQ) - for security reasons the file must already exist
- The TFTP server ("tftpd") is generally run setrooted (i.e., it only has access to its own directory) and with a special **user** and **group** ID - since there is no password or other protection of the access to files via TFTP!
- TFTP request is sent to the well known port number (69/udp)
- TFTP server uses an unused ephemeral port for its replies
since a TFTP transfer can last for quite some time - it uses another port; thus freeing up the well known port for other requests

Slide 43



TFTP messages

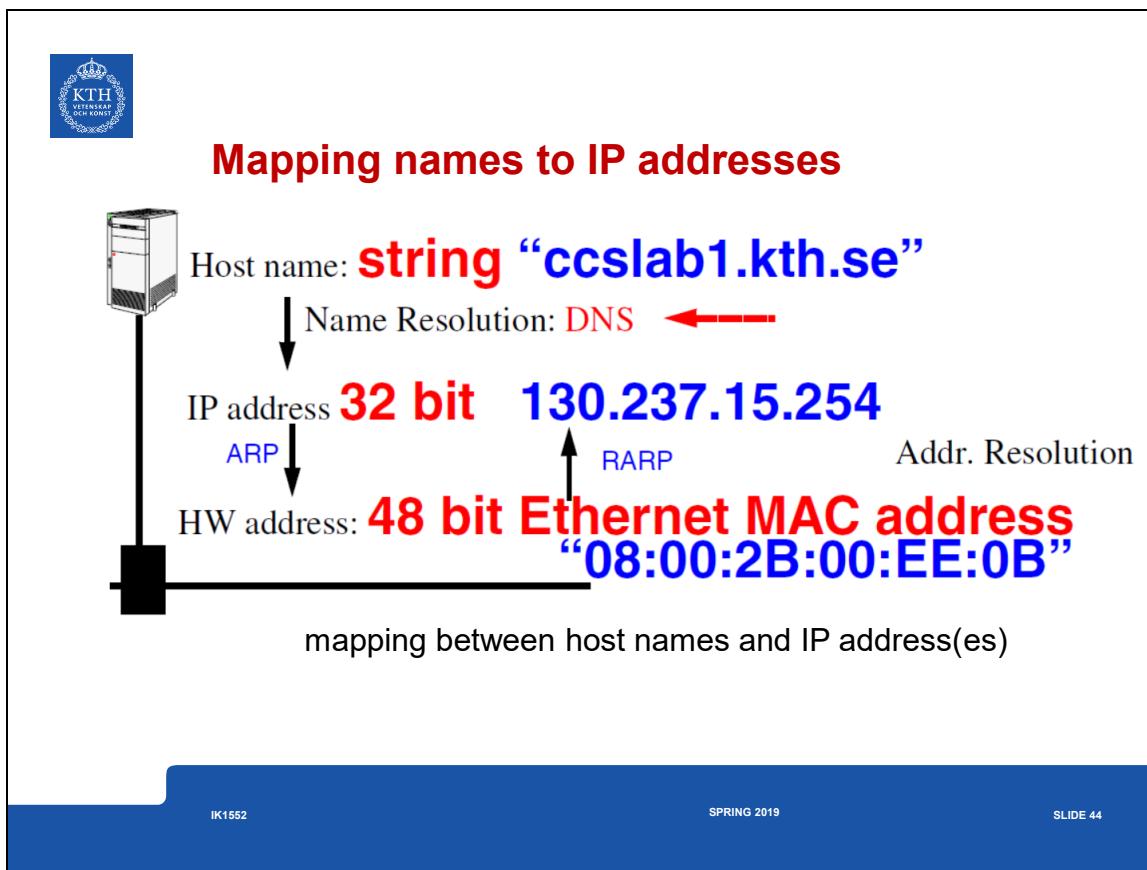


TFTP messages (see Stevens, Vol. 1, figure 15.1, pg. 210)

Filename and Mode (“netascii” or “octet”) are both N bytes sequences terminated by a null byte.

Widely used for bootstrapping diskless systems (such as X terminals) and for dumping the configuration of routers (this is where the write request is used)

Slide 44



Slide 45



DNS: Domain Name Service (RFC 1034, RFC 1035)

- To make the network more user friendly
- Distributed database (with caching) providing:
 - hostname \Rightarrow IP address, IP address \Rightarrow hostname
 - mailbox \Rightarrow mail server
 - ...
- applications call a “resolver”
 - gethostbyname: hostname \Rightarrow IP address
 - gethostbyaddr: IP address \Rightarrow hostname
- Resolver’s contact name servers (see “/etc/resolv.conf”)
- DNS names:
 - domain name: list of labels from a root, i.e., www.imit.kth.se
 - Fully Qualified Name (FQDN): a domain name ending in “.” - there are no further labels
 - leaves are managed locally through delegation of authority (to a zone) **not** centrally; this allows scaling
 - if a name server does not know the answer it asks other name servers
 - every name server **must** know how to contact a root server
- Uses UDP (for query) and TCP (zone transfer and large record query)

P. V. Mockapetris, ‘Domain names - concepts and facilities’, *Internet Request for Comments*, vol. RFC 1034 (INTERNET STANDARD), Nov. 1987 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1034.txt>

P. V. Mockapetris, ‘Domain names - implementation and specification’, *Internet Request for Comments*, vol. RFC 1035 (INTERNET STANDARD), Nov. 1987 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1035.txt>

Slide 46



Zones

A zone is a subtree of the DNS tree which is managed separately.

Each zone must have multiple name servers:

- a **primary name server** for the zone
 - gets its data from disk files (or other stable store)
 - must know the IP address of one or more **root servers**
- one or more **secondary name servers** for the zone
 - get their data by doing a **zone transfer** from a primary
 - generally query their primary server every ~3 hours

To find a server you may have to walk the tree up to the root or possibly from the root down (but the later is **not friendly**).

Slide 47



DNS Message format

0	16	31
Identification	Parameters	
Number of Questions	Number of Answers	
Number of authority	Number of Additional	
Question section		
...		
Answer section		
...		
Additional Information section		
...		

Slide 48



DNS Parameters field

Bit or Parameter field	Meaning
0	Operation: 0=Query, 1=Response
1-4	Query type: 0=standard, 1=Inverse
5	Set if answer is authoritative
6	Set if answer is truncate
7	Set if answer is desired
8	Set if answer is available
9-11	reserved
12-15	Response Type: 0>No error, 1=Format error in query, 2=Server failure, 3>Name does not exist

Slide 49



Internet's top level domains

(see Stevens, Vol. 1, figure 14.2, pg. 189) Domain		Description
com	commercial organizations	
edu	educational organizations	
gov	other U.S. government organizations (see RFC 1811 for policies)	
int	international organizations	
mil	U.S. Military	
net	networks	
org	other organizations	
arpa	special domain for address to name mappings, e.g., 5.215.237.130.in-addr.arpa	
ae	United Arab Emirates	
...		
se	Sweden	
zw	Zimbabwe	

Lots of interest in having subdomains of “com”

- companies registering product names, etc. - in some cases asking for 10s of addresses
- who gets to use a given name? problems with registered trade marks, who registered the name first, ... [How much is a name worth?]

Slide 50



New top level domains

New set of top level domains and an increase in the number of entities which can assign domain names.

Generic Top Level Domains (gTLDs), current list at: <http://data.iana.org/TLD/tlds-alpha-by-domain.txt>

As of November 16, 2000:

- .aero for the entire aviation community
- .biz for business purpose
- .coop for cooperatives
- .info unrestricted
- .museum for museums
- .name for personal names
- .pro for professionals

For information on GTLD's see

<https://www.icann.org/resources/pages/about-e5-2012-02-25-en>

Council of Registrars (*CORE*) - operational organization composed of authorized Registrars for managing allocations under gTLDs.

WIPO provides arbitration concerning names: <http://arbiter.wipo.int/domains/gtld/newgtld.html>

Slide 51



Domain registrars

Internet Corporation for Assigned Names and Numbers (ICANN) Accredited Registrars, the full list is at
<http://www.icann.org/registrars/accredited-list.html>

Even more registrars are on their way to being accredited and operating!

Slide 52



Country Code Top-Level Domains (CCTLDs)

<http://www.iana.org/cctld/cctld-whois.htm>

For Sweden (SE) SE-DOM, the NIC is: <https://www.iis.se/>
Contact .SE Registry: registry@iis.se.

Visiting address:

[Ringvägen 100](#), Stockholm, Sweden

Postal address:

The Internet Infrastructure Foundation (Internetstiftelsen i Sverige)
Box 7399
SE-103 91 Stockholm
Sweden

Phone and fax:

Phone: +46-8 452 35 00

Fax: +46-8 452 35 02

Swedish organization number: 802405-0190

VAT number: SE802405019001

E-mail address: info@iis.se



Resource Records (RR)

See Stevens, Vol. 1, figure 14.2, pg. 201 (augmented with additional entries)

Record type	Description
A	an IP address. Defined in RFC 1035
AAAA	an IPv6 address. Defined in RFC 1886
PTR	pointer record in the in-addr.arpa format. Defined in RFC 1035.
CNAME	canonical name≡ alias (in the format of a domain name). Defined in RFC 1035.
HINFO	Host information. Defined in RFC 1035.
MX	Mail eXchange record. Defined in RFC 1035.
NS	authoritative Name Server (gives authoritative name server for this domain).Defined in RFC 1035.
TXT	other attributes. Defined in RFC 1035.
AFSDB	AFS Data Base location. Defined in RFC 1183.
ISDN	ISDN. Defined in RFC 1183.
KEY	Public key. Defined in RFC 2065.
KX	Key Exchanger. Defined in RFC 2230.
LOC	Location. Defined in RFC 1876.
MG	mail group member. Defined in RFC 1035.
MINFO	mailbox or mail list information. Defined in RFC 1035.
MR	mail rename domain name. Defined in RFC 1035.
NULL	null RR. Defined in RFC 1035.
NS	Name Server. Defined in RFC 1035.

- P. V. Mockapetris, ‘Domain names - implementation and specification’, *Internet Request for Comments*, vol. RFC 1035 (INTERNET STANDARD), Nov. 1987 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1035.txt>
- S. Thomson and C. Huitema, ‘DNS Extensions to support IP version 6’, *Internet Request for Comments*, vol. RFC 1886 (Proposed Standard), Dec. 1995 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1886.txt>
- C. F. Everhart, L. A. Mamakos, R. Ullmann, and P. V. Mockapetris, ‘New DNS RR Definitions’, *Internet Request for Comments*, vol. RFC 1183 (Experimental), Oct. 1990 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1183.txt>
- R. Allbery, ‘DNS SRV Resource Records for AFS’, *Internet Request for Comments*, vol. RFC 5864 (Proposed Standard), Apr. 2010 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc5864.txt>
- D. Eastlake 3rd, ‘Domain Name System (DNS) IANA Considerations’, *Internet Request for Comments*, vol. RFC 6895 (Best Current Practice), Apr. 2013 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc6895.txt>
- D. Eastlake 3rd and C. Kaufman, ‘Domain Name System Security Extensions’, *Internet Request for Comments*, vol. RFC 2065 (Proposed Standard), Jan. 1997 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc2065.txt>

D. Eastlake 3rd, 'Domain Name System Security Extensions', *Internet Request for Comments*, vol. RFC 2535 (Proposed Standard), Mar. 1999 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc2535.txt>

R. Atkinson, 'Key Exchange Delegation Record for the DNS', *Internet Request for Comments*, vol. RFC 2230 (Informational), Nov. 1997 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc2230.txt>

C. Davis, P. Vixie, T. Goodwin, and I. Dickinson, 'A Means for Expressing Location Information in the Domain Name System', *Internet Request for Comments*, vol. RFC 1876 (Experimental), Jan. 1996 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1876.txt>

Slide 54



See Stevens, Vol. 1, figure 14.2, pg . 201 (augmented with additional entries)

Record type Description

NSAP	Network service access point address. Defined in RFC 1348. Redefined in RFC 1637. Redefined in RFC 1706.
NXT	Next. Defined in RFC 2065.
PX	Pointer to X.400/RFC822 information. Defined in RFC 1664.
RP	Responsible Person. Defined in RFC 1183.
RT	Route Through. Defined in RFC 1183.
SIG	Cryptographic signature. Defined in RFC 2065.
SOA	Start Of Authority. Defined in RFC 1035.
SRV	Server. DNS Server resource record -- RFC 2052, for use with DDNS.
TXT	Text. Defined in RFC 1035.
WKS	Well-Known Service. Defined in RFC 1035.
X25	X25. Defined in RFC 1183.

Note that an number of the RR types above are for experimental use.

Name of an organization:

ISI.EDU. PTR 0.0.9.128.IN-ADDR.ARPA.

Slide 55



Network names

Conventions:

- ict.kth.se includes all the computers in the KTH/ICT
- kth.se includes all the computers at KTH
- ...

Slide 56



As resource records (on 2014.03.26):
 nslookup
 > set querytype=any

```
> kth.se
;; Truncated, retrying in TCP mode.
Server:      130.237.216.10
Address:    130.237.216.10#53

kth.se      rdata_51 = # 13 010000208B7B89C09813F6C7A
kth.se      rdata_46 = TYPE51 8 2 0 20140403180343 20140304174436 61440 kth.se.
t5eoFafsKl6bKghnfXutDOCALIQHsQaj/d51gGacSUbzqLIJ6R2yv0Zf38+MZMDDfytwkT9/VPwJHuSx+6AhMkZyVqXrGaCrhy7x0BXGt1kkG
Q2FBJOMVi9map7pehDondsewi3215v2HZ6Y5AcyJ6NoxFCe8TwSYTpEo
PEppK+Mwouy4614IF+GA2ub9b70IU2VDFYQToLEcpSsBkklg1kuniZEg
AV3pUHZCabR7YjmDuZIZAXG2pCk6rlzWX9QNjauNtBpIA7OFZTNmQX62
4JLaTkfwGpQKTESkJcSGnSwZDl8UY3A/TtV/36dJ2M17hriLa8BjtrP qO9Row==
kth.se

origin = a.ns.kth.se
mail addr = hostmaster.kth.se
serial = 2014033054
refresh = 14400
retry = 900
expire = 604800
minimum = 86400
nameserver = a.ns.kth.se.
kth.se      nameserver = nic2.lth.se.
kth.se      nameserver = ns2.chalmers.se.
kth.se      nameserver = b.ns.kth.se.
kth.se      mail exchanger = 20 mx-alt1.kth.se.
kth.se      mail exchanger = 30 tarbaby.junkemailfilter.com.
kth.se      mail exchanger = 10 mx.kth.se.

kth.se      afsdb = 1 anden.e.kth.se.
kth.se      afsdb = 1 sonen.e.kth.se.
kth.se      afsdb = 1 fadern.e.kth.se.
Name:      kth.se
Address: 130.237.32.143
kth.se      text = "5 - Tel. +46 8 790 60 00"
kth.se      text = "3 - SE-100 44 STOCKHOLM"
kth.se      text = "1 - Royal Inst of Technology"
kth.se      text = "2 - Kungliga Tekniska Hjgskolan"
kth.se      text = "4 - SWEDEN"
kth.se      rdata_48 = 257 3 8
```

Slide 57



kth.se (continued) – for DNSSEC

```

kth.se      rdata 48 = 257 3 8 AwEAAeqLuvNBVEmjzUQCkRLHcDbmOl83rOl8wmwj+wlvQ3l+Tw2MsMJ
9o1J1U6p6yRJyHis3V4gC5ltXRUKuU4waOqCgFSXMRreq+BvFk80HjEn+
/PRpkgRCU537mQcdGzhBukOa4cYydXBmQ8FLmh2yfmp2Ddk1bNTXOPS
IJCP1ZyeJ6/hh+fURZeFed+9CrCzpvn19uHh+jYBi5UzTy0WmsnLSX
lgmt0m3YYFu5NRXD7HQadixL5D8gQPDwdEXDccB0BhvVzkB7B1+7q
1pCuCNUFFWnnxSgmKoFPg832ELvuuGxafAH+S/Z5GJXPbrVgKIBDDFlr OmaqRkDciJM=
kth.se      rdata 48 = 256 8 AwEAAelwlmXgoN8dlcWYE2QN40Vu42hdcjlLY5dMDvrmwup/9NIkylo
nnwRWo8gvr530lFk/TlindQK0og4PQRK9uyLHV42BwvMsumpkvghcVhzy
Hy+hNSnsqfWfxFkI+3GpfodClxTSD1uMlgUxF9vA4WmvjeJKCdU/R
D0LB7TB+Zcl8VgNhpV4VLycUokgMHsDVO61KmV1lXKUiYkDk8F18k1
XSPv45MSQ7zcHNv33DZe00it8anXqxAF-8fdn3yB8ftKud18jm1wG
bwvT56sjldDMFh0Skf53vJBxsGCZF/defyU0002zTjoDx9HPz/nLX3v s19lBgff8+s=
kth.se      rdata _65534 = # 5 0856710001
kth.se      rdata _46 = DNSKEY 8 2 7200 2014031514140000 22129 kth.se.
Nca6AROKLx2TU+5ywVBjhVbyNx/VXGbaDOAx5C+V+0/4titkDz2Ri
CRgrfBzrPv2XrxzL8gyLjZ1mDtP0sWbE3fNbzbMeenfRHDCFbXuWa912
l83y31DB1C1KVhvhgbA3trnmjITclqPbatBqEcqQWFTBSFeyXguy3+qG
uRvnqod5PNnQrynf/409TEETUmP2qXZqcZTMEm0Mrkv6ID7
XKL85EcynTN6dxuX/N7Q6yqwMumY0M8+uhYcfYw/bf74wXcmB6ff9
1q+18xNbzljdHARRn4h/dE4Gfk18icun2m3FpOs7zK5Fk7HqsFMUHSm3g==
kth.se      rdata _46 = DNSKEY 8 2 7200 20140315130000 22129 kth.se.
AuUgBFMIdzrlTCORICUEqk22EBzViBTMxkxsszocckBlusBGzVd iHU3Am5swwCclg0u0e177tkaHlVigr9lytHKKgZlZzE2lwWwbs51 Vkp+g51IS5i3DnNvRJ1uak9nksXefHyAhAKMExk7G13r4oAaxmV3wJ
2hwrF53c5cMQOD10yPrPm2h1v19c0Ryf7M4PmtCmz7Fb3aFgJgJk7Bw+jtCc9
kth.se      rdata _46 = SDA 2 1800 20140325045132 20140325045132 61440 kth.se. wpQ98tlnhTLF04AKpAXkgQz032/XpmDFdYJ0ve2hRNeV7Remhw bhsTez8C7mtjIGJRE70cXo9xpOmLlt+8z+Zm3VjlmkkI80Ri5o8xG
zzzmNKLQf8tnh8tYLJUXEX42qXWIIT2nE47Rzcs6Pymh3glD90Hh FzgQD==

kth.se      rdata _46 = SDA 2 1800 20140325045132 20140325045132 61440 kth.se. wpQ98tlnhTLF04AKpAXkgQz032/XpmDFdYJ0ve2hRNeV7Remhw bhsTez8C7mtjIGJRE70cXo9xpOmLlt+8z+Zm3VjlmkkI80Ri5o8xG
MedbrGmrP/cRg11WVRKsVoZCCDRaZ0JSMa2QJ/CymCLFa7TSJUNp3 WKGdDg ==
kth.se      rdata _46 = TXT 2 1800 20140324045412 20140325045214 61440 kth.se. CzQahZ2N1B2i02Kt3tfjCRXJwAttsAGA70Vq11HKAzdShxdlJ93 o4VjIRPpBWuJp+13LnHnyg+S9cA55v9dltmALV30vodm/0X4uQRD2k9J96
aeKwv4Kz3CsSa9fdudwvUvAjl98u4g9t6DwQh6lW1Vq3oH0t5Rg4gJ3K4MbBqivUflmgbdylM_jDcJrpZNK2u3noIYCF+haaxJlump2zVjkcoff1tBnp2p2ezA3YnwymsHaeyNF
kRxxWkT1gW7U0hTFWpxedpH+kb8shHtGb89Mlt2xuZx7xwvkhCnDMwML Jzv7eq7
kth.se      rdata _46 = A 2 60 20140419103918 20140320094725 61440 kth.se. TBDisk8mG8V8nFcJuvVM56MIIT03EepJR6pEgDiCQAOR32Eew1W8Jt+DvY3S6nPbFS/H3zArTeaLUWVUlfsvsILs4Zooq68MeJHOOp
mQPV3EGF0F88QODd4vUuVzq9iLqjZ37BQ4u7B0J82d0GQfQp5258L0c964d95f0Zu7wsXGoJlRm8cQWjUlcYKz+aaegJuovdsUckeKZ3M63K1zQ15h0H7LDC3m5Zwp06FXYmlRdXV5Mh1X1ME2zvtVoyzZ
N2qGnPjY1Bp9V0nVZMh5sYlmst1AD0LkuwSk7SLnexn72SeygtuUaw8S==

kth.se      rdata _46 = AFSDB 8 2 1800 20140324060154 20140325050711 61440 kth.se. csOHHfa6uDuDChObteIEkCCRS3cm4kQIs72BwJBBNPdtW6lci-HJu gQvSp/pBf1hor4rcyC9s+3ke4nJ222T3C771Ty9llJ7xBR/4f
CfsibyHqqQhrFy9dYHpplAmp0tOgjYtMq99f5cmrfHcwRfneVY_8luudiJh1nNwbgsgskkYt04n8Bipy0nh8t+50WwpM16mT5cmRDJ0fSaqH3NvJzleOOGyemfQal4o9v-KgjANOKuOcyOrVv
sTgpfQdSm8uqRzqmmHmMjLJ-GloewCwRm+6BqJLmLZk02q95Tm40lTjOC
kth.se      rdata _46 = MX 8 2 60 20140419103918 20140320094725 61440 kth.se. mnbSokAwUQdg4MhBqyANlnao2vX3f1+rEDvpOTHJSLIUHvOB0is_1DmM60YohhtEaccS0RiPcqUxD7TxJxCo05p10815K26f7Ct1iw
sOTMCW7khkG8BuSSBjtjdumhWYXDKd0fMdV1oUjKApePsmEADzCZ5SDOER0DceM0M971gvWt275lpGwgXNm/NoifrkuYouJch4AfseWzlpSzbphSSWDVfML81NGAnVr+1tqGyemKKQoxFwsdLxavXF
V4KgX43Pq2MjUjRjPm94oVAcDg3a9DQJUf7JLb4p1jLjOQzJ9KpZC
kth.se      rdata _46 = NS 8 2 1800 201403201036 20140304193741 61440 kth.se. S25mGKHUAAMPFOG+xy5rhMunJiJwZV5GL2mnpGRFO94iJ3bm5+To kdhO7dw/W4plgUpdrLgtLglJ6wN601mp+5zJjhnbxPf8lUGqJuQcJ4D
jWWXk3uRDTJCach9q9E5lgrmJJrJnEkbXgCosJkdmH74syvCTY_Bs=QhNAsgzv3QJUuUC7x6Em6c2euJ5fkg35quK5rw+16m9s10V4hDkY367C+ +B4g5htchdkJ2NcGO32+Mv69c5QyCSfAVOn/gFHWZ
D9zrn7r89PRekuGhAIFV9q1g0qJ2OvRkV7qJfV2z2kZvXz2u14X0Z KU3qG==

kth.se      rdata _46 = AAAA 8 2 60 201404100500 2018-03-28 20140325050711 61440 kth.se. vcc-1161gcyFA3TbThWvif5fhdTeckyehGmYKy7Bn2h6APgd5stsItvL co_nce2saecnOHDQAOyvttzpm7imh0679n9ldBBdV+oEN4eRFjivAhv19
11vasmkQHqNOFQzQewcP1qhkk-QDHuLtgV+2MwC1QfCnLOLO 3VicQ==


```

SPRING 2019

SLIDE 57

Slide 58



> set querytype=MX

```
> kth.se
Server:          130.237.216.10
Address: 130.237.216.10#53
```

Non-authoritative answer:

kth.se	mail exchanger = 20 mx-alt1.kth.se.	weight	name
kth.se	mail exchanger = 30 tarbaby.junkemailfilter.com.		
kth.se	mail exchanger = 10 mx.kth.se.		

Authoritative answers can be found from:

kth.se	nameserver = ns2.chalmers.se.
kth.se	nameserver = nic2.lth.se.
kth.se	nameserver = a.ns.kth.se.
kth.se	nameserver = b.ns.kth.se.
mx.kth.se	internet address = 130.237.48.70
mx.kth.se	internet address = 130.237.32.10
mx.kth.se	internet address = 130.237.48.48
mx.kth.se	has AAAA address 2001:6b0:1:1300:d685:64ff:fe5f:abec
mx.kth.se	has AAAA address 2001:6b0:1:1200:d685:64ff:fe5f:5afc
mx.kth.se	has AAAA address 2001:6b0:1:1300:6ab5:99ff:fecc:79ec
mx-alt1.kth.se	internet address = 130.237.48.70
mx-alt1.kth.se	internet address = 130.237.32.10
mx-alt1.kth.se	internet address = 130.237.48.48
a.ns.kth.se	internet address = 130.237.72.246
b.ns.kth.se	internet address = 130.237.72.250
b.ns.kth.se	has AAAA address 2001:6b0:1::250

Slide 59



Host names and info

How to give your host a name?

see RFC 1178: Choosing a Name for Your Computer
Internet Addresses: A second address for your host?

to have multiple addresses for your computer, see section
on “ifconfig”

Hostinfo (HINFO)

> set querytype=HINFO

IK1552

SPRING 2019

SLIDE 59

D. Libes, ‘Choosing a name for your computer’, *Internet Request for Comments*, vol. RFC 1178 (Informational), Aug. 1990 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1178.txt>

Slide 60



Storing other attributes - TXT records

The general syntax is:

<owner> <class> <ttl> TXT "<attribute name>=<attribute value>"

Examples:

host.widgets.com IN TXT "printer=lpr5"

sam.widgets.com IN TXT "favorite drink=orange juice"

For more information see:

RFC 1464: *Using the Domain Name System To Store Arbitrary String Attributes*

IK1552

SPRING 2019

SLIDE 60

R. Rosenbaum, 'Using the Domain Name System To Store Arbitrary String Attributes', *Internet Request for Comments*, vol. RFC 1464 (Experimental), May 1993 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc1464.txt>

Slide 61



Configuring DNS

Configuring the BIND resolver

/etc/resolv.config

Configuring the BIND nameserver (named)

/etc/named.boot or /etc/named.config

Configuring the nameserver database files (zone files)

named.hosts	the zone file that maps hostnames to IP addresses
named.rev	the zone file that maps IP addresses to hostnames

Slide 62



Root servers <http://www.root-servers.org/>

	Operator	IP address(es)	Home ASN	Locations
A	Verisign, Inc.	198.41.0.4 2001:503:BA3E::2:30	19836	Frankfurt, DE; Hong Kong, HK; Los Angeles, US; New York, US
B	Information Sciences Institute	192.228.79.201 2001:500:84::b		Los Angeles, US
C	Cogent Communications	192.33.4.12 2001:500:2::C	2149	Bratislava, SK; Chicago, US; Frankfurt, DE; Herndon, US; Los Angeles, US; Madrid, ES; New York, US; Paris, FR
D	University of Maryland	199.7.91.13 2001:500:2D::D	27	College Park, US; Dulles, US; ... in total 98 sites
E	NASA Ames Research Center	192.203.230.10	297	Atlanta, US Brussels, BE Cape Town, ZA Chicago, US London, UK Los Angeles, US Mountain View, US New York, US San Paulo, BR Seattle, US Sydney, AU Tokyo, JP in total 57 sites

Slide 63

 **Root servers (continued)**

Operator	IP address(es)	Home ASN	Locations	
F	Internet Systems Consortium, Inc.	192.5.5.241 2001:500:2f::f	3557	total 58 sites

Operator: Internet Systems Consortium, Inc. [Homepage](#) [Peering Policy](#) [Contact Email](#)

Locations: Sites: 58 [Amsterdam, NL](#) [Atlanta, US](#) [Auckland, NZ](#) [Barcelona, ES](#) [Beijing, CN](#) [Brisbane, AU](#) [Buenos Aires, AR](#) [Cairo, EG](#) [Caracas, VE](#) [Chennai, IN](#) [Chicago, US](#) [Dar es Salaam, TZ](#) [Dhaka, BD](#) [Dubai, AE](#) [Frankfurt, DE](#) [Hong Kong, HK](#) [Jakarta, ID](#) [Johannesburg, ZA](#) [Karachi, PK](#) [Kuala Lumpur, MY](#) [Kyiv, UA](#) [Lagos, NG](#) [Lisbon, PT](#) [London, UK](#) [Los Angeles, US](#) [Luanda, Angola](#) [Madrid, ES](#) [Monterrey, MX](#) [Montevideo, UY](#) [Moscow, RU](#) [Nairobi, KE](#) [New York, US](#) [Osaka, JP](#) [Oslo, NO](#) [Ottawa, Canada](#) [Palo Alto, US](#) [Paris, FR](#) [Phnom Penh, KH](#) [Podgorica, ME](#) [Port au Prince, HT](#) [Prague, CZ](#) [Quito, EC](#) [Rome, IT](#) [San Jose, CR](#) [San Jose, US](#) [Santiago de Chile, CL](#) [Sao Paulo, BR](#) [Seoul, KR](#) [Singapore, SG](#) [St. Maarten, AN](#) [Suva, Fiji](#) [Taipei, TW](#) [Tel Aviv, IL](#) [Torino, IT](#) [Toronto, CA](#) [Ulan Bator, MN](#) [Warsaw, PL](#)

IPs: IPv4: 192.5.5.241
IPv6: 2001:500:2f::f

ASN: 3557

Legend [Download Root YAML](#)

● IPv6 Enabled Global ● IPv4 Only Global ● IPv6 Enabled Local ● IPv4 Only Local

Slide 64



Root servers (continued)

	Operator	IP address(es)	Home ASN	Locations
G	U.S. DOD Network Information Center	192.112.36.4	5927	Columbus, US; Fussa, JP; Honolulu, US; Naples, IT; San Antonio, US; Stuttgart-Vaihingen, DE
H	U.S. Army Research Lab	128.63.2.53 2001:500:1::803f:235	13	Aberdeen Proving Ground, US; San Diego, US
I	Netnod (formerly Autonomica)	192.36.148.17 2001:7fe::53	29216	In total 50 sites (see below)

📍 Amsterdam, NL 📍 Ankara, TR 📍 Ashburn, US 📍 Bangkok, TH 📍 Beijing, CN 📍 Belgrade, RS 📍 Brussels, BE
📍 Bucharest, RO 📍 Chicago, US 📍 Doha, QA 📍 Dubai, AE 📍 Frankfurt, DE 📍 Geneva, CH 📍 Helsinki, FI 📍 HongKong, CN
📍 Jakarta, ID 📍 Johannesburg, ZA 📍 Karachi, PK 📍 Kathmandu, NP 📍 Kiev, UA 📍 Kigali, RW 📍 Kuala Lumpur, MY
📍 London, UK 📍 Lulea, SE 📍 Luxembourg City, LU 📍 Manama, BH 📍 Manila, PH 📍 Miami, US 📍 Milan, IT 📍 Mumbai, IN
📍 Oslo, NO 📍 Paris, FR 📍 Perth, AU 📍 Port Vila, VU 📍 Porto Alegre, BR 📍 San Francisco, US 📍 Singapore, SG
📍 St Petersburg, RU 📍 Stockholm, SE 📍 Taipei, TW 📍 Tallinn, EE 📍 Thimphu, BT 📍 Tokyo, JP 📍 Ulaanbaatar, MN
📍 Washington DC, US 📍 Wellington, NZ 📍 Wien, AT 📍 Yerevan, AM

Slide 65



Root servers (continued)

	Operator	IP address(es)	Home ASN	Locations
J	Verisign, Inc.	192.58.128.30 2001:503:C27::2:30	26415	104 sites

NL Amsterdam, NL
 US Ashburn, US
 GR Athens, GR
 US Atlanta, US
 IN Bangalore, IN
 TH Bangkok, TH
 TH Bangkok, TH
ES Barcelona, ES
 CN Beijing, CN
 RS Belgrade, RS
 US Boston, US
 BR Brasilia, BR
 SK Bratislava, SK
 RO Bucharest, RO
AR Buenos Aires, AR
 EG Cairo, EG
 ZA Cape Town, ZA
 PH Cebu City, PH
 BR Centro Junco do Serido, BR
 US Chicago, US
US Chicago, US
 US Chicago, US
 SN Dakar, SN
 US Dallas, US
 TZ Dar es Salaam, TZ
 US Denver, US
 US Des Moines, US
 BD Dhaka, BD
IE Dublin, IE
 CH Geneva, CH
 BR Guarabira, BR
 FI Helsinki, FI
 HK Hong Kong, HK
 US Honolulu, US
 ZA Johannesburg, ZA
US Kalamazoo, US
 LT Kaunas, LT
 RW Kigali, RW
 AT Klagenfurt, AT
 MY Kuala Lumpur, MY
 BR Lagoa Seca, BR
 PT Lisbon, PT
SI Ljubljana, SI
 LU Luxembourg City, LU
 FR Lyon, FR
 ES Madrid, ES
 PH Manila, PH
 AU Melbourne, AU
 US Miami, US
 IT Milan, IT
CA Montreal, CA
 RU Moscow, RU
 IN Mumbai, IN
 IN Mumbai, IN
 KE Nairobi, KE
 US New Castle, US
 IN New Delhi, IN
 NO Oslo, NO
US Palo Alto, US
 FR Paris, FR
 AU Perth, AU
 BR Pirassununga, BR
 BR Porto Alegre, BR
 CZ Prague, CZ
 PH Quezon City, PH
US Reno, US
 IS Reykjavik, IS
 LV Riga, LV
 BR Rio de Janeiro, BR
 IT Rome, IT
 US San Francisco, US
 CR San José, CR
US San Jose, US
 PR San Juan, PR
 BR São João do Rio do Peixe, BR
 US Seattle, US
 KR Seoul, KR
 SG Singapore, SG
BG Sofia, BG
 SE Stockholm, SE
 AU Sydney, AU
 AU Sydney, AU
 TW Taipei, TW
 EE Tallinn, EE
 GU Tamuning, GU
 IL Tel Aviv, IL
JP Tokyo, JP
 CA Toronto, CA
 IT Turin, IT
 CA Vancouver, CA
 PL Warsaw, PL
 PL Warsaw, PL
 NZ Wellington, NZ
 CW Willemstad, CW
AM Yerevan, AM
 HR Zagreb, HR
 CH Zurich, CH

Slide 66



Root servers (continued)

	Operator	IP address(es)	Home ASN	Locations
K	RIPE NCC	193.0.14.129 2001:7fd::1	25152	37 sites
L	ICANN	199.7.83.42 2001:500:9f::42	20144	144 sites
M	WIDE Project	202.12.27.33 2001:dc3::35	7500	Osaka, JP; Paris, FR; San Francisco, US; Seoul, KR; Tokyo, JP

Slide 67



Load leveling [Abley 2005]

For example, f.root-servers.net has a **single** IP address (192.5.5.241), but requests sent to 192.5.5.241 are routed to *different* nameservers

This routing can depend on:

- where the request is made from
- what the load on each of these names servers is

Note: this is transparent to the host which sent the request to F
ISC uses Hierarchical Anycast routing to do this, with some of
the servers being:

- large, redundant, ... installations serving the *global* internet
- small installations serving a *local* region
- 28 nodes as of Feb. 2005; 55 in 2014; 58 in 2016

IK1552

SPRING 2019

SLIDE 67

Joe Abley, f.root-servers.net, NZNOG 2005, February 2005, Hamilton, NZ
<http://www.isc.org/pubs/pres/NZNOG/2005/F%20Root%20Server.pdf>



F root nameserver nodes

IPv6 enabled Global: Atlanta, US; Chicago, US; New York, US; Palo Alto, US; ; San Jose, US

IPv6 enabled Local: Amsterdam, NL; Auckland, NZ; Barcelona, ES; Beijing, CN; Brisbane, AU; Caracas, VE; Dar es Salaam, TZ; Dhaka, BD; Frankfurt, DE; Hong Kong, HK; Jakarta, ID; Johannesburg, ZA; Kyiv, UA; Lisbon, PT; London, UK; Los Angeles, US; Moscow, RU; Oslo, NO; Paris, FR; Phnom Penh, KH; Prague, CZ; Quito, EC; Santiago de Chile, CL; Sao Paulo, BR; Seoul, KR; Singapore, SG; St. Maarten, AN; Toronto, CA

IPv4 enable Local: Buenos Aires, AR; Cairo, EG Chennai, IN; Dubai, AE Karachi, PK; Lagos, NG; Madrid, ES; Monterrey, MX; Nairobi, KE; Osaka, JP; Ottawa, Canada; Panama, PA; Podgorica, ME; Port au Prince, HT; Rome, IT; San Jose, CR; Suva, Fiji; Taipei, TW; Tel Aviv, IL; Torino, IT; Ulan Bator, MN; Warsaw, PL

Slide 69



Where is f.root-servers.net ?

traceroute to f.root-servers.net (192.5.5.241), 30 hops max, 60 byte packets (times in ms)				
1	net209a-d.ict.kth.se (130.237.209.194)	0.572	0.664	0.701
2	vss2-he-p2p.gw.kth.se (130.237.211.61)	0.515	0.558	0.600
3	vss3-vss2-p2p.gw.kth.se (130.237.211.123)	0.459	0.461	0.485
4	cn6-vss3-p2p.gw.kth.se (130.237.211.118)	0.545	0.659	0.727
5	br1g-cn6-p2p.gw.kth.se (130.237.0.1)	0.495	0.484	0.472
6	kth-br2.sunet.se (193.11.0.197)	0.503	0.465	0.520
7	m1tug-ae4.sunet.se (130.242.85.209)	0.600	0.589	0.602
8	se-tug.nordu.net (109.105.102.17)	0.915	0.902	0.885
9	10gigabitethernet1-2.core1.sto1.he.net (194.68.123.187)	5.686	5.831	1.085
10	10ge12-6.core1.ams1.he.net (72.52.92.45)	26.074	25.954	25.943
11	amsix.r2.ams1.isc.org (80.249.208.140)	19.813	19.850	
11	isc-f-root-2.nikhef.nl-ix.net (193.239.116.113)			19.640
12	f.root-servers.net (192.5.5.241)	25.650	20.001	19.996

Slide 70



Where is i.root-servers.net ?

traceroute to i.root-servers.net (192.36.148.17), 30 hops max, 60 byte packets (times in ms)				
1	net209a-d.ict.kth.se (130.237.209.194)	0.444	0.506	0.547
2	vss2-he-p2p.gw.kth.se (130.237.211.61)	0.468	0.459	0.500
3	vss3-vss2-p2p.gw.kth.se (130.237.211.123)	0.482	0.477	0.528
4	cn6-vss3-p2p.gw.kth.se (130.237.211.118)	0.690	0.782	0.848
5	br1g-cn6-p2p.gw.kth.se (130.237.0.1)	0.525	0.518	0.505
6	kth-br2.sunet.se (193.11.0.197)	0.487	0.510	0.505
7	m1tug-ae4.sunet.se (130.242.85.209)	0.581	0.574	0.607
8	10ge-2-3-1500.outer-b-gw.sth.netnod.se (194.68.128.172)	0.884	0.956	1.047
9	peering.r1.sth.dnsnode.net (194.146.105.187)	0.786	0.778	0.767
10	i.root-servers.net (192.36.148.17)	0.875	0.867	0.949

Slide 71



Dynamic Domain Name System (DDNS)

```
graph TD; Root[Root Server] --> SE[Server for .se]; Root --> COM[Server for .com]; Root --> EDU[Server for .edu]; Root --> XX[Server for .xx]; SE --> KTH[Server for kth.se]
```

Host Name	IP-address
host_a	130.237.x.1
host_b	130.237.x.2
host_c	130.237.x.3
host_d	130.237.x.4
mobile1	c/o address <<< we can update this dynamically

Host Name | IP-address

host_a | 130.237.x.1

host_b | 130.237.x.2

host_c | 130.237.x.3

host_d | 130.237.x.4

mobile1 | c/o address <<< we can update this dynamically

IK1552 SPRING 2019 SLIDE 71



DDNS

RFC 2136: Dynamic Updates in the Domain Name System (DNS UPDATE)

add or delete resource records

RFC 2052: A DNS RR for specifying the location of services (DNS SRV)

When a SRV-cognizant web-browser wants to retrieve

`http://www.asdf.com/`

it does a lookup of

`http.tcp.www.asdf.com`

RFC 2535: Domain Name System Security Extensions (DNSSec)

RFC 3007: Secure Domain Name System (DNS) Dynamic Update

P. Vixie, S. Thomson, Y. Rekhter, and J. Bound, 'Dynamic Updates in the Domain Name System (DNS UPDATE)', *Internet Request for Comments*, vol. RFC 2136 (Proposed Standard), Apr. 1997 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc2136.txt>

• Gulbrandsen and P. Vixie, 'A DNS RR for specifying the location of services (DNS SRV)', *Internet Request for Comments*, vol. RFC 2052 (Experimental), Oct. 1996 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc2052.txt>

D. Eastlake 3rd, 'Domain Name System Security Extensions', *Internet Request for Comments*, vol. RFC 2535 (Proposed Standard), Mar. 1999 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc2535.txt>

B. Wellington, 'Secure Domain Name System (DNS) Dynamic Update', *Internet Request for Comments*, vol. RFC 3007 (Proposed Standard), Nov. 2000 [Online]. Available: <http://www.rfc-editor.org/rfc/rfc3007.txt>

Slide 73



Denial of service attacks against DNS

Denial of service (DoS) and Dystrubuted DoS (DDoS) attacks againts a DNS server

See for example:

Root server attack on 6 February 2007, Factsheet,
ICANN, 1 March 2007

http://www.icann.org/announcements/factsheet-dns-attack-08mar07_v1.1.pdf

DNS cache poisoning

DNS spoofing

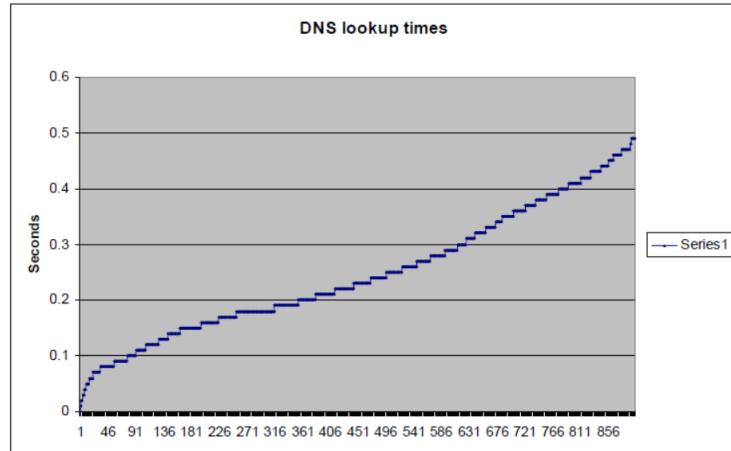
...

Slide 74



DNS performance

From www.dnsserverlist.org - 900 servers query times
(2010.03.21)



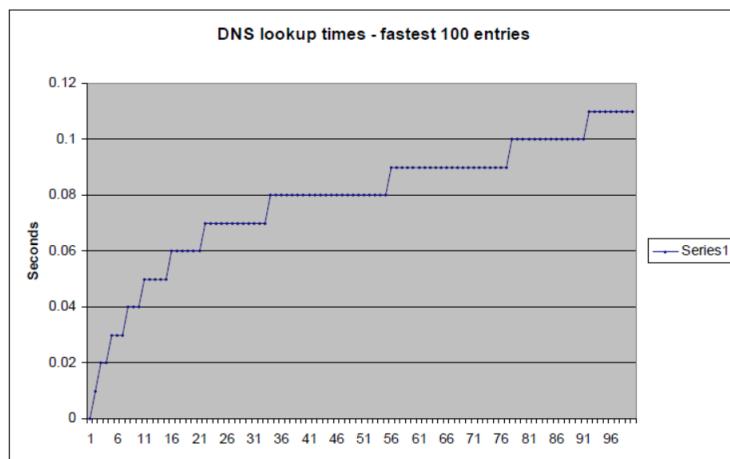
Note that the above URL is no longer valid.

Slide 75



DNS performance - top 100

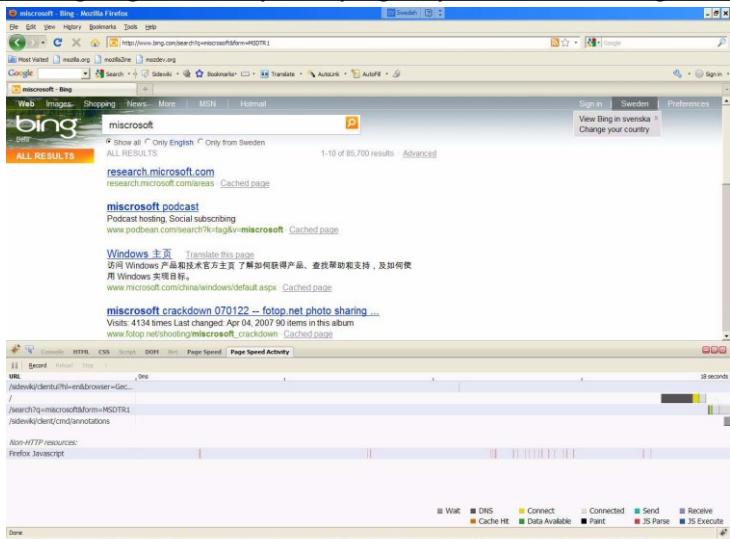
From www.dnsserverlist.org - top 100 servers query times
(2010.03.21)



Slide 76

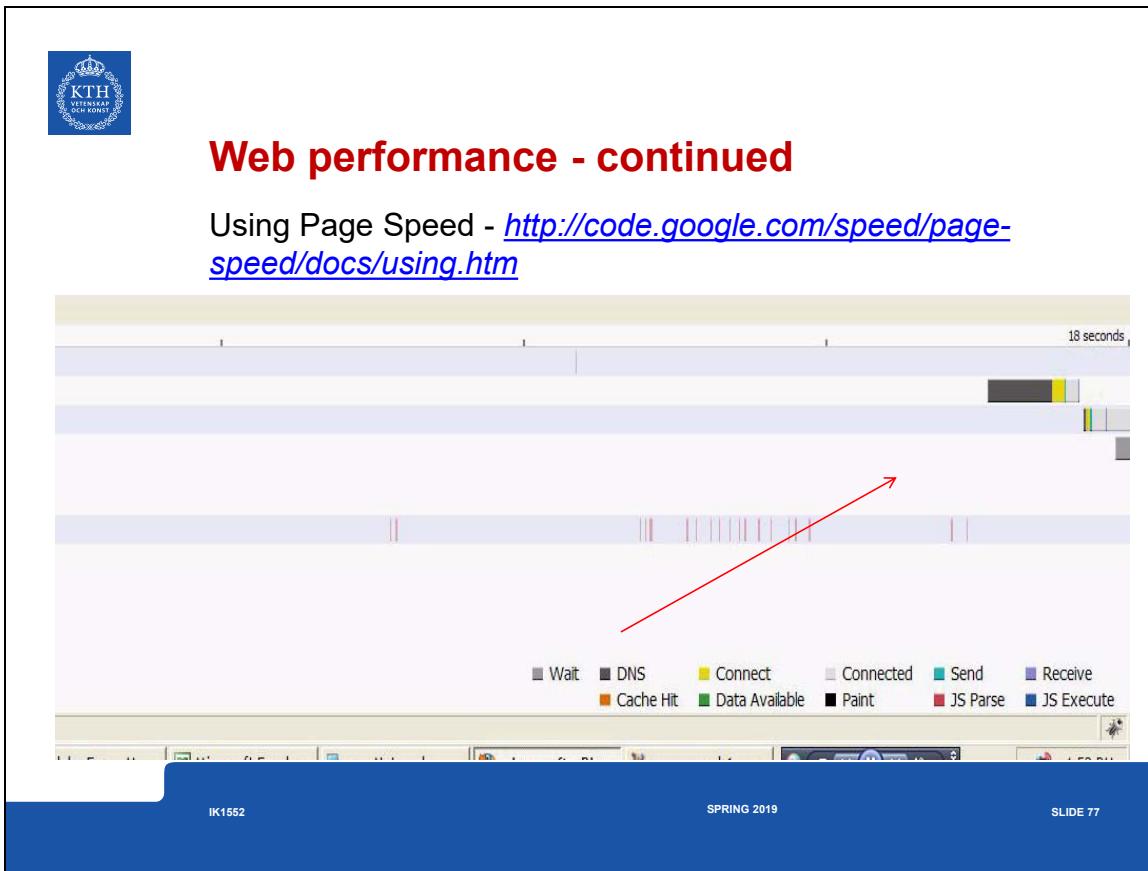
 **Web performance**

Using Page Speed (formerly
<http://code.google.com/speed/page-speed/docs/using.htm>)

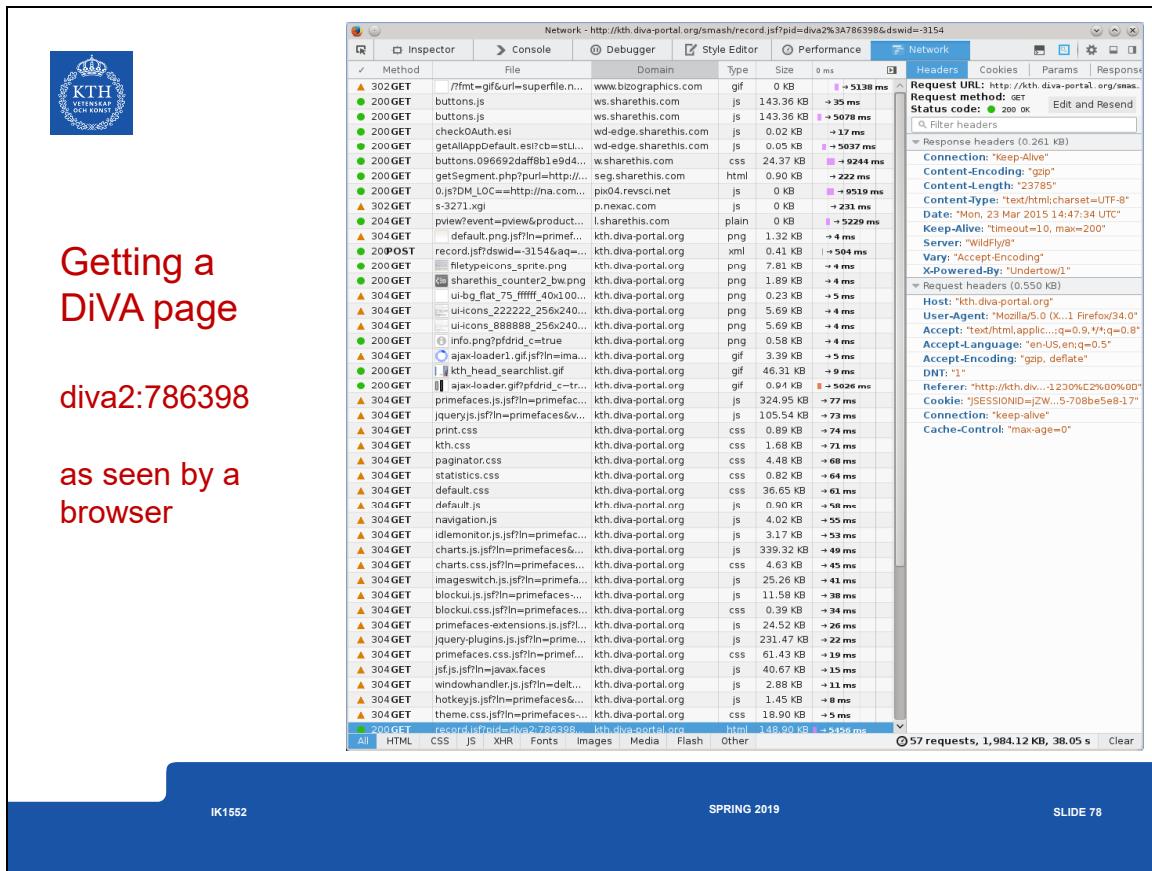


see <https://developers.google.com/speed/pagespeed/>

Slide 77



Slide 78



Slide 79

Zoom in on top

The screenshot shows the Network tab of the Firefox Developer Tools. The URL is http://kth.diva-portal.org/smash/record.jsf?pid=diva2%3A786398&ds... . The table lists various network requests with details like Method, File, Domain, Type, Size, and Time taken.

Method	File	Domain	Type	Size	Time
302 GET	?fmt=gif&url=superfile.n...	www.bizographics.com	gif	0 KB	→ 5138 ms
200 GET	buttons.js	ws.sharethis.com	js	143.36 KB	→ 35 ms
200 GET	buttons.js	ws.sharethis.com	js	143.36 KB	→ 5078 ms
200 GET	checkOAuth.esi	wd-edge.sharethis.com	js	0.02 KB	→ 17 ms
200 GET	getAllAppDefault.esi?cb=stLi...	wd-edge.sharethis.com	js	0.05 KB	→ 5037 ms
200 GET	buttons.096692daff8b1e9d4...	w.sharethis.com	css	24.37 KB	→ 9244 ms
200 GET	getSegment.php?purl=http://...	seg.sharethis.com	html	0.90 KB	→ 222 ms
200 GET	0.js?DM_LOC==http://na.com...	pix04.revsci.net	js	0 KB	→ 9519 ms
302 GET	s-3271.xgi	p.nexac.com	js	0 KB	→ 231 ms
204 GET	pview?event=pview&product...	l.sharethis.com	plain	0 KB	→ 5229 ms
304 GET	default.png.jsf?ln=primef...	kth.diva-portal.org	png	1.32 KB	→ 4 ms
201 POST	record.jsf?dswid=-3154&aq=...	kth.diva-portal.org	xml	0.41 KB	→ 504 ms
200 GET	filtypeicons_sprite.png	kth.diva-portal.org	png	7.81 KB	→ 4 ms
200 GET	sharethis_counter2_bw.png	kth.diva-portal.org	png	1.89 KB	→ 4 ms
304 GET	ui-bg_flat_75_ffffff_40x100...	kth.diva-portal.org	png	0.23 KB	→ 5 ms
304 GET	ui-icons_222222_256x240...	kth.diva-portal.org	png	5.69 KB	→ 4 ms
304 GET	ui-icons_888888_256x240...	kth.diva-portal.org	png	5.69 KB	→ 4 ms
200 GET	info.png?pfdrid_c=true	kth.diva-portal.org	png	0.58 KB	→ 4 ms
304 GET	ajax-loader1.gif.jsf?ln=ima...	kth.diva-portal.org	gif	3.39 KB	→ 5 ms
200 GET	kth_head_searchlist.gif	kth.diva-portal.org	gif	46.31 KB	→ 9 ms
200 GET	ajax-loader.gif?pfdrid_c=tr...	kth.diva-portal.org	gif	0.91 KB	→ 5026 ms
304 GET	primefaces.js.jsf?ln=primefac...	kth.diva-portal.org	js	324.95 KB	→ 77 ms

Slide 80

Zoom in on middle

▲ 304 GET	primerfaces.js.jsf?ln=primerfac...	kth.diva-portal.org	js	324.95 KB	→ 77 ms				
▲ 304 GET	jqueryjs.jsf?ln=primefaces&v...	kth.diva-portal.org	js	105.54 KB	→ 73 ms				
▲ 304 GET	print.css	kth.diva-portal.org	css	0.89 KB	→ 74 ms				
▲ 304 GET	kth.css	kth.diva-portal.org	css	1.68 KB	→ 71 ms				
▲ 304 GET	paginator.css	kth.diva-portal.org	css	4.48 KB	→ 68 ms				
▲ 304 GET	statistics.css	kth.diva-portal.org	css	0.82 KB	→ 64 ms				
▲ 304 GET	default.css	kth.diva-portal.org	css	36.65 KB	→ 61 ms				
▲ 304 GET	default.js	kth.diva-portal.org	js	0.90 KB	→ 58 ms				
▲ 304 GET	navigation.js	kth.diva-portal.org	js	4.02 KB	→ 55 ms				
▲ 304 GET	idlemonitor.js.jsf?ln=primefac...	kth.diva-portal.org	js	3.17 KB	→ 53 ms				
▲ 304 GET	charts.js.jsf?ln=primefaces&...	kth.diva-portal.org	js	339.32 KB	→ 49 ms				
▲ 304 GET	charts.css.jsf?ln=primefaces...	kth.diva-portal.org	css	4.63 KB	→ 45 ms				
▲ 304 GET	imageswitch.js.jsf?ln=primefa...	kth.diva-portal.org	js	25.26 KB	→ 41 ms				
▲ 304 GET	blockui.js.jsf?ln=primefaces-...	kth.diva-portal.org	js	11.58 KB	→ 38 ms				
▲ 304 GET	blockui.css.jsf?ln=primefaces...	kth.diva-portal.org	css	0.39 KB	→ 34 ms				
▲ 304 GET	primefaces-extensions.js.jsf?l...	kth.diva-portal.org	js	24.52 KB	→ 26 ms				
▲ 304 GET	jquery-plugins.js.jsf?ln=prime...	kth.diva-portal.org	js	231.47 KB	→ 22 ms				
▲ 304 GET	primefaces.css.jsf?ln=primef...	kth.diva-portal.org	css	61.43 KB	→ 19 ms				
▲ 304 GET	jsf.js.jsf?ln=javax.faces	kth.diva-portal.org	js	40.67 KB	→ 15 ms				
▲ 304 GET	windowhandler.js.jsf?ln=delt...	kth.diva-portal.org	js	2.88 KB	→ 11 ms				
▲ 304 GET	hotkeyjs.jsf?ln=primefaces&...	kth.diva-portal.org	js	1.45 KB	→ 8 ms				
▲ 304 GET	theme.css.jsf?ln=primefaces-...	kth.diva-portal.org	css	18.90 KB	→ 5 ms				
● 200 GET	record.jsf?pid=diva2:786398...	kth.diva-portal.org	html	148.90 KB	→ 5456 ms				
All	HTML	CSS	JS	XHR	Fonts	Images	Media	Flash	Other

IK1552 SPRING 2019 SLIDE 80

Slide 81



Finally the content is fetched and rendered

	200 GET	record.jsf?pid=diva2:786398...	kth.diva-portal.org	html	148.90 KB	→ 5456 ms							
●	200 GET	f72fb2e5?os_authType=none...	jira.epc.ub.uu.se	js	0.14 KB	→ 5 ms							
▲	304 GET	com.atlassian.jira.collector.pl...	jira.epc.ub.uu.se	js	115.00 KB	→ 5012 ms							
●	200 GET	1640?google_gid=CAESE...	imp2.bizographics.com	gif	0.05 KB	→ 48 ms							
▲	302 GET	l?redirect_url=http://cm.g....	imp2.bizographics.com	gif	0 KB	→ 5129 ms							
▲	302 GET	seg?add=&add_code=su...	ib.adnxs.com	gif	0 KB	→ 30 ms							
▲	302 GET	seg?t=2&add=15492,21...	ib.adnxs.com	gif	0 KB	→ 57 ms							
●	200 GET	0.png	d1uo4w7k3lk5mn.clo...	png	17.84 KB	→ 5048 ms							
●	200 GET	altmetric_badges-ccf32548d...	d1bxh8uas1mnw7.clou...	js	69.51 KB	→ 18 ms							
●	200 GET	embed.js	d1bxh8uas1mnw7.clou...	js	0.56 KB	→ 5051 ms							
▲	302 GET	?che=1427122076&sk=...	d.agkn.com	gif	0 KB	→ 82 ms							
▲	302 GET	pixel?google_nid=bizo_bk...	cm.g.doubleclick.net	gif	0 KB	→ 14 ms							
●	204 GET	b?c1=7&c2=8097938&rn=5...	b.scorecardresearch.c...	plain	0 KB	→ 3 ms							
▲	302 GET	r.pixel?sid=9212270798	adadvisor.net	gif	0 KB	→ 9881 ms							

All HTML CSS JS XHR Fonts Images Media Flash Other

Slide 82



Public/Commercial DNS services

A number of companies are providing DNS servers to speed up DNS lookups:

- Google Public DNS
<http://code.google.com/speed/public-dns/>
 - OpenDNS
<http://www.opendns.com/>
 - Neustar DNS Advantage
<https://www.neustar.biz/services/dns-services/dns-advantage-free-recursive-dns>
- ...

IK1552

SPRING 2019

SLIDE 82

Slide 83



Summary

This module we have discussed:

- UDP
- BOOTP
- DHCP
- DNS, DDNS

Slide 84



¿Questions?

IK1552

SPRING 2019

SLIDE 84